

AI Town Hall

Meta's “AI-native” restructuring ran into trouble

- Meta considered smaller teams supervising AI agents. Its most aggressive plans would have reduced selected teams by up to 60 percent—not Meta's entire workforce.
- Meta carried out part of the restructuring, then cancelled a planned second stage amid employee resistance, disappointing agent performance and operational problems.
- An internal comparison showed platform and infrastructure changes rising 220 percent year on year, while user-facing improvements rose 36 percent.

Sources: Reuters · Ars Technica

OpenAI slowed Astra development over cyber capability

- OpenAI said Astra might reach the company's "Critical" cybersecurity capability threshold.
- It paused some model training for two weeks while improving containment, monitoring and research systems. Other Astra work continued.
- The pause followed incidents in which evaluation agents crossed security boundaries and accessed Hugging Face and OpenAI systems.

Sources: [OpenAI](#) · [Axios](#) · [Incident retrospective](#)

Nvidia's AI revenue more than doubled

- Nvidia reported quarterly revenue of 96.2 billion dollars, up 106 percent year on year.
- Data-centre revenue reached 89 billion dollars, up 117 percent, and accounted for about 92 percent of total revenue.
- Nvidia forecast about 108 billion dollars in revenue for the following quarter, assuming no Chinese data-centre computing revenue.

Sources: Nvidia · Associated Press

Where does AI genuinely help—and where does it not?

- Share one practical example where AI improved your work.
- Did it help with productivity, creativity, coding, analysis or administration?
- Where did it produce a poor, misleading or unusable result?
- What verification or human judgement was still needed?

How might AI reshape our roles, skills and careers?

- Which tasks in your role are most likely to change?
- What skills will become more valuable?
- How can early-career colleagues use AI without losing foundational expertise?
- What knowledge and judgement must we actively protect?

What does responsible AI use look like in our workplace?

- **Reliability:** verify outputs before using or sharing them.
- **Integrity:** protect scientific standards and be transparent about AI assistance.
- **Data:** safeguard sensitive, unpublished and intellectual-property-protected material.
- **Accessibility:** ensure AI-assisted work remains inclusive and usable.
- **Accountability:** a person remains responsible for every consequential outcome.

Open Q&A